


RADBOD UNIVERSITY NIJMEGEN



FACULTY OF SOCIAL SCIENCE

A title

A PROJECT FOR THE AGES

THESIS MSc ARTIFICIAL INTELLIGENCE

Student Information

Surname: HOMMELSHEIM
First name: RON
Student number: s1000522
E-mail address: RON.HOMMELSHEIM@RU.NL
Course code: SOW-MK194 (45 EC)
AI Specialisation: COGNITIVE COMPUTING

Supervisor Information

Role: SUPERVISOR 1
Surname: THILL
First name: SERGE
Institute: RADBOUD UNIVERSITY
E-mail address: SERGE.THILL@DONDEERS.RU.NL
Supervision type: INTERNAL

Abstract

Contents

1	Problem Statement	3
2	Background	3
2.1	Program Synthesis	3
3	DreamCoder	3
3.1	DreamCoder	3
4	DeepSynth	4
5	Search space	4
6	Program space	4
7	GFlowNet	6
8	Theory	6
8.1	Introduction	6
9	Methods	7
9.1	Complexity analysis	8

1 Problem Statement

It has long been hypothesized that our cognition is based on some sort of language of thought (LOT). Recent advances in the field have shown promising results by interpreting this LOT as a type of programming language [1, 2, 3]. The language is constructed from the bottom-up using primitives in a domain specific language. However, thus far this has been done purely syntactically. I hypothesize that meaning is not just an emergent phenomenon but a crucial part of concept formation. I therefore propose using graph neural networks and represent program components as nodes, in order to include relational structure. Additionally, nodes can contain supplementary semantic information. This hypothesis will be tested by comparing the resulting model to Ellis et al.’s DreamCoder [2].

Let’s assume the premise that the mind is computational, in one form or another. and we are building a world model which can be thought of as probabilistic programs.

2 Background

2.1 Program Synthesis

3 DreamCoder

One of the most fascinating undertakings of Artificial Intelligence (AI) is its objective to reveal and learn from the mechanisms of human cognition.

The idea that the thoughts are constructed by some kind of Language of Thought (LOT), in order to understand and represent our reality has a long history. Gottfried Leibniz imagined a *Characteristica Universalis*, a formal language capable of expressing metaphysical concepts [4]. Jerry Fodor, among many other scholars, developed this hypothesis, and proposed that thoughts are composed according to some form of internal grammar [5]. If our thoughts are indeed composed by some language, we can let our intuitions about its structure and limitations be guided by the advances made in the past century in regards to formal languages and computationalism. Gödel showed that any formal system strong enough to express Peano Arithmetic is incomplete [6]. Turing showed that any possible computation can be done by a Turing machine and concluded that the human mind must be computational [7]. [Piantadosi disagrees] Chomsky’s Hierarchy classifies formal languages of different strengths and the automata that recognize them [8]. These results and many more are consequential when thinking about a language of thought. Many important questions about concepts are still left unanswered, e.g. regarding their formation, the structure of the framework they reside in, their representation, how they connect to other concepts and so on. Many contem-

porary versions of the LOT hypothesis (LOTH) argue that our conceptual framework is constructed bottom-up using some initial primitives. More so, the language we construct concepts in is executable, akin to a programming language [3]. The idea is that we model the world by creating programs that generate our observations [1]. Roumi et al.’s experiment indeed suggest that humans may interpret and compress sequence regularities in a type of program of minimum description length (MDL) [9]. Dehaene et al. further test this hypothesis and find similar results [3]. The question then remains, where primitives actually come from. Piantadosi asserts that symbols have no inherent meaning, instead, meaning emerges from the dynamic relations between symbols. I.e., only the conceptual role and the dynamics of the symbols define meaning [10]. By defining meaning as an emergent phenomenon of conceptual roles, Piantadosi and Hill do not rule out that large language models (LLM) already have some foundation of meaningful concepts [11]. Recent work tries to relate these ideas to biological systems, e.g. Quan et al. propose how role-filler interactions could be implemented in the brain via temporal synchrony to permit symbolic processing and dynamic binding of variables [12]. Santoro et al. argue for a semiotic approach, assert that symbols are subjective, and propose that meaning of symbols in artificial systems will arise when they are subjected to socio-cultural interactions [13].

Ellis et al. use λ -calculus, which is Turing complete, to build something akin to a LOT [2]. In response, Piantadosi shows that combinatorial logic is equivalent to λ -calculus and why combinatorial-logic-like (LCL) language is preferable, for one because it doesn’t assume primitives, thereby avoiding the problem of their meaning. One of the questions regarding a LOT is whether thoughts can be produced by simple syntactic symbol manipulation. I (and many others) suspect that semantics are not just an emergent phenomenon but that it actively shapes our perception and concept formation [13, 14]. Moreover, purely symbolic attempts, albeit being the original approach of artificial intelligence, have proven to be insufficient [15]. Instead, the successes of modern AI are owed to the advances of connectionist models. These however, have their own deficits, namely, being incredibly data-hungry, lacking causality, and lacking out-of-distribution generalization, to name a few. Therefore, new approaches try to combine these two AI strands into what is known as *neurosymbolic AI* [15].

3.1 DreamCoder

DreamCoder (DC) is a model that synthesises programs from initial primitives in a domain specific language (DSL) and a set of tasks [2]. The model utilises a modified version of a wake-sleep algorithm [16] to learn a generative model and a recognition network in tandem. The generative model learns a probability distribution over

programs while the recognition network learns to map from tasks to programs and serves as neurally guided search over the program space. DreamCoder builds its library of programs from primitives in its DSL, and creates higher-order functions. Sub-routines that are commonly used, are chunked and abstracted into concepts which become more accessible. This narrows the search tree immensely and enables scalability. Tasks can be generative (e.g. creating an image) or conditional, (input-output relation, e.g. sorting a list).

The model operates in three distinct phases.

Wake In the wake phase, the objective is to find the best program ρ for each task $x \in X$ DreamCoder is presented with. Each task consist of one or a few (up to ~ 10) examples. The system synthesises the programs using a neural network as the recognition model Q to search through its library D .

Sleep: Abstraction A key component of the success of DreamCoder is the refactoring of programs. In this phase, DreamCoder consolidates new abstractions from the programs it has found during the wake phase. Experiences are replayed, and commonalities between programs are found so that they can be abstracted into new concepts to use in the future. The objective is to find programs of minimum description length (MDL).

Sleep: Dreaming The objective during the dreaming phase is to train the recognition model Q to perform MAP inference by both fantasies, and replays. When fantasising, programs are drawn from the learned library as a generative model. The model is then trained on the produced tasks. This ensures a large and varied data-set to learn from and encourages out-of-distribution generalisation. Replays, simply training on previously seen task-program pairs, ensures that the model improves its predictions on the actual tasks it needs to learn, by using its programs.

A Bayesian View - What exactly happens? - Include the parameterization of the generative model as a bigram transition matrix (three dimensional. Parent, arg index, depth?)

4 DeepSynth

- Introducing different search strategies. The two main approaches are enumerative search and sampling. (also show entropy sampling (kilcher's explanation)) However, either way, they are first predicting weights for the PCFG and then searching or sampling from the PCFG.

5 Search space

6 Program space

- How can we calculate the size of that space? - Hierarchical structure - Abstract Syntax Trees

Recent advancements in program synthesis have seen the rise of novel techniques aimed at generating code that not only satisfies a given specification but also matches human-written code patterns. One of the most notable methods in this sphere is DreamCoder. This model, along with several other state-of-the-art (SOTA) methodologies, has shown exceptional proficiency in understanding and generating complex code structures based on given specifications.

However, a critical examination of such methods, including DreamCoder, reveals a foundational limitation: their heavy reliance on syntactical constraints. While these constraints are undoubtedly vital for ensuring the correctness of generated programs, they do not necessarily guarantee a deep understanding or utilization of semantic relationships within the code. This is akin to learning the grammar of a language but missing out on the nuances and contexts that give meaning to words and phrases.

One could argue that the attention mechanism, as seen in transformers, provides a potential solution to this limitation. The attention mechanism inherently captures dependencies between various parts of an input, no matter how distant they might be. In the context of program synthesis, such a mechanism could be instrumental in deriving and maintaining meaningful semantic relationships between different code segments. The beauty of attention is that it provides a dynamic way to weigh different parts of the input based on context, potentially allowing a transformer-based model to capture intricate semantic details of a program.

Nevertheless, a transformer's capabilities come with its own set of challenges. While the attention mechanism offers a promising avenue for capturing semantic nuances, transformers, especially when trained from scratch, require colossal amounts of data to achieve satisfactory performance. The fundamental architecture of transformers necessitates this data-driven learning, often demanding diverse and extensive training datasets that might not always be available or feasible, especially in niche areas of program synthesis. Moreover, when the goal is to create a model of human cognitive abilities, it is obvious that humans do not have access to vast amounts of data, instead, we are able to infer causal relationships, from only a limited number of examples. E.g. consider the problem $[1,2,3] \rightarrow [3,2,1]$. An astute reader might promptly conclude that the list has been reversed. Dougs analogies

This is where GFlowNets (GFNs) carve their niche. Contrary to the data-hungry nature of transformers,




Figure 1: (A) 8 different domains of tasks. (B) Representation of the learned library of concepts. On the left we see the initial primitives from which the concepts in the middle region are composed. On the right we see a task as input-output relations and the found solution. On the bottom is the same solution expressed only with initial primitives. Image taken with permission from the original paper [2].

GFNs do not necessitate gargantuan datasets for efficient training. Instead, they leverage a more intuitive approach: defining a reward distribution that’s directly proportional to the objects (or programs) we intend to generate. This not only streamlines the learning process but also ensures that the synthesized programs are aligned with our desired outcomes.

In the realm of program synthesis, our objectives are clearly defined. We aim to construct objects, or more specifically, pieces of code that precisely meet the given specification. This isn’t just about producing syntactically correct code; it’s about generating semantically rich and efficient programs that can effectively solve the specified problem without superfluous or redundant components.

7 GFlowNet

The problem can be formalized as finding a latent hierarchical structure from a limited set of specifications.

We can specify our GFlowNet in the realm of program synthesis as building a directed acyclic graph over abstract syntax trees. We formalise it. The gfn is conditioned on the task, and so we have a conditional reward distribution $R(z|x)$, as well as a conditional forward policy $P_F(s|x)$ and partition function $Z(x)$.

I am building on the DeepSynth framework where the goal is to find programs in a domain specific language.

The task is list editing, i.e. i get some inputs and outputs relationship and the network has to find a program that solves it.

we want diversity, so perhaps even multiple programs that solve it.

8 Theory

The trained GFlowNet gives us a the stochastic policy $\pi(a|s)$, where a is an action from the action space A and s is a state from the state space S .

- Relation to Reinforcement learning - Relation to MCMC sampling - Relation to Variational inference.

Since their initial publication [SOURCE], many extended and modified variants have been published. See e.g. [SOURCE, awesome GFlowNets for an overview.]

8.1 Introduction

A Generative Flow Network, or GFlowNet, operates as a generative model driven by a trained stochastic policy. It constructs objects $z \in Z$, where Z is the space of completed states, sequentially, with each sampling probability being proportional to a reward function $R(z)$, where $R(z)$ is non-negative and integrable. GFlowNets excel in sampling diverse solutions, eliciting a high reward [17].

The state space can be visualized as a directed acyclic graph (DAG), where vertices correspond to states and edges denote transitions.

A trajectory τ is a series of state transitions commencing from an initial state s_0 and culminating at a terminal state, $s_n \in Z$.

The forward policy, denoted as $P_F(s'|s)$, encompasses the children of all non-terminal states in $S \setminus Z$. It inherently generates a distribution over complete trajectories:

$$P(\tau) = \prod_{i=1}^n P_F(s_i|s_{i-1})$$

Using sequential sampling from P_F , one can deduce a distribution P_F^\top over terminal states:

$$P_F^\top(z) = \sum_{\tau \text{ leading to } z} P_F(\tau)$$

For any reward function $R : Z \rightarrow \mathbb{R}_{\geq 0}$, GFlowNets aims to determine a parametric policy where object sampling likelihood is proportional to its reward.

The main theorem describes that if the flow function F is trained such that it matches the flow-matching constraint, i.e. for any state the flow going into the state matches the flow going out of it, and the flow at terminated states x is defined by the reward function $R(x)$, GFlowNet will sample terminated states with probability $\frac{R(x)}{\sum_{x'} R(x')}$.

Since its original publication many extended and modified variants have been published.

1. **TB Objective:** Trajectory Balance (TB) approach, as elaborated by Malkin et al. (2022), necessitates simultaneous learning of a forward policy, a backward policy $P_B(s|s';\theta)$, and a scalar Z_θ . The TB objective is:

$$L_{TB}(\tau; \theta) = \left[\log \frac{Z_\theta \prod_{i=1}^n P_F(s_i|s_{i-1}; \theta)}{R(z) \prod_{i=1}^n P_B(s_{i-1}|s_i; \theta)} \right]^2$$

However, since I am essentially predicting a linearized tree, each node can only have one parent. Therefore, P_B will always be 1 and can be disregarded from the equation. We can then simplify the TB Objective, making use of log rules, to

$$L_{TB}(\tau; \theta) = \left(\log Z_\theta + \sum_{i=1}^n \log P_F(s_i|s_{i-1}; \theta) - \log R(z) \right)^2$$

2. **Training Dynamics:** Nullifying L_{TB} ensures proportional sampling to object reward. The loss can be minimized using gradient descent with on-policy and off-policy strategies, reminiscent of reinforcement learning techniques.

9 Methods

In a probabilistic programming representation of world models and thoughts, as a type of Language of Thought, which aligns with constructivism we can conceive of the problem statement as

In a language of thought, regardless of most details, we tend to think in a paradigm in which thoughts are compositional. In one way or another, thoughts are hierarchical (?). Both in ontologies, i.e. the way concepts are structured, (animal - bird - fink) but also in the sequential nature of thought construction (as in natural language, reasoning tasks, etc.). We are creating parse trees, or abstract syntax trees.

The renowned professor Yoshua Bengio explains that thoughts are composed of multiple intermediate steps and explains that reasoning, or thinking, is the combination of knowledge and inference. He theorises that we we construct a world model which may take the form of an energy based model. It is interesting to think about whether this world model is explicit or implicit. In DreamCoder and other SOTA methods in program synthesis the world model is essentially an explicit PCFG from which one samples programs. In neural program synthesis, the new approach of neurosymbolic programming, the role of the neural network is to act as a recognition model, helping to search programs from this PCFG. I imagine it a little different, and see the thought construction process more like an agent, following a policy.

On one side, in alignment with Occam's razor, our world model aims to identify the essential knowledge fragments required to interpret observed data. This model should be concise, optimizing the encoding bits and enhancing the reutilization of knowledge segments, thereby enhancing its generalization capability to unfamiliar scenarios. This aligns with the principle of parsimoniousness or, as Dehaene describes as finding programs of minimum description length.

A trained inference machine amortizes the cost of searching through a vast solution space which resource-intensive search methods like MCMC, traditional AI techniques, or optimization algorithms cannot handle. In conclusion, approximate inference could be the solution here, since a larger model doesn't lead to overfitting but better approximations, while not being dependent on lots of data.

Yoshua warns that a notable vulnerability of current large language models is the conflation of the world model with the inference mechanism. Although LLMs may construct an implicit world model, the distinction between world model and inference may be the underlying cause of their difficulty for out-of-distribution prediction [18, 19].

In his book, *Thinking Fast and Slow*, Nobel-prize laureate Daniel Kahneman popularized the notion of two systems being the core mechanisms of human cognition [SOURCE]. - explain system 1 and two - explain hemi-

spheric lateralization, mcgilchrist, etc. - notion that if we don't have a concept for it, it does not exist for us. I.e. even if our sensory equipment can in principle take it in, and does take it in, it is not there for us. It seems we are unable to attend to it. Perhaps attribute meaning to it and therefore it doesn't enter our conscious awareness. - I think meaning is the integrated framework of relationships between concepts. This is the field we observe the world through. - Concepts in culture and society, i.e. over different levels in a hierarchy. (Memes) - Computational boundary of the self. How do multicellular agents, or multiple agents compute information together? There must be a common language, that is approximately the same. - Allostasis & Homeostasis (of concepts/ memes) - Imaginaria, where do unicorns exist, etc. What is a table - Can any concept be definitively defined as a particular concept (either symbolic or symbolic vector e.g. SPA, hyperdimensional vectors) or are all concepts approximations? - Introduce different frameworks and architectures (Kissner, Japa, Concepts (energy based model from i think deepmind), GLOM, SPA, DreamCoder) - Introduce GFlowNets as a way to create thought trajectories - Explain that GFlowNets could be used over multiple of these ideas and frameworks (discrete, continuous, etc.) it is a general framework. - Here I'm going to attempt to use it for probabilistic program synthesis.

Humans naturally exhibit the aptitude to innately evaluate the reliability of our knowledge. Understanding the limits of one's knowledge directs the acquisition of new information, thus optimizing the learning process.

Moreover, human cognition is characterized by modularity and reusability. Our mental architectures are adept at disentangling knowledge into modular components, which can be recalled, combined, and repurposed to address novel challenges. By continuously crafting and refining abstract concepts, humans can process, assimilate, and communicate multifaceted knowledge structures with remarkable efficiency.

Further, we intuitively discern causal relationships, allowing for predictive and counterfactual reasoning. Causal models, seen as extensive distribution families influenced by interventions, provide the scaffolding for this intuitive understanding. As highlighted by Kemp et al. (2015), these models enable humans to extrapolate beyond observed data, envisage hypothetical scenarios, and navigate the world with a goal-oriented mindset. In computational terms, integrating causality can tremendously amplify the generalization capabilities of algorithms, allowing them to predict outcomes under novel interventions and facilitating deeper understanding of intricate systems.

Often there are multiple solutions to a query, which is why the inference network should model a multimodal distribution rather than a point estimate. Difference between MAP and MLE.

Stephen Wolfram explains in his book about ChatGPT

and the accompanying blog article how a thought can be seen as a trajectory through semantic space. - introduce transformers, how they work and especially the attention mechanism. - introduce the notion of spaces

This also relates to causality.

In a probabilistic programming paradigm, we view this as compositional functions, with the many properties of functional programming analogous to currying, etc.

What is a thought? [Thoughts as trajectories] Perhaps we can think of understanding, as integrating information into a meaningful whole that is largely consistent.

We can formalise this as a hierarchical latent structure Z .

Framing the problem.

Combinatorial search problem.

- If we are using encoder + decoder, are we violating the Markovian Flow assumption? Look at the smiley example. If the NN would get [[left brow], [left brow, right brow], [left brow, right brow, smile]] as input, i.e. the sequence of the states, it would violate the assumption. but it only gets the current state, e.g. [left brow, right brow], and from that it has to infer the next step. When using a decoder, we would indeed give it the whole trajectory of states, so it would violate the assumption. But even now i am encoding the sequence and giving the whole trajectory as input. and since the CFG is essentially a tree, there is only one parent for each state.

It would probably be faster to do it bottom up like in HEAP search from Nathanael and from GFN-EM, then we could also use sub-trajectory balance, i.e. calculate intermediate Rewards.

Another thing we could do is predict a bunch of terminals at once, and then combine in each step.

Other ideas: wave function collapse

- Reasons for using a GFlowNet

- Transformer allows for semantic relationships (although we still don't use control or data flow, but we could in principle)

- We want to approximate a multimodal distribution, unlike DreamCoder in which the MAP estimate refrains from finding semantically equal but syntactically different solutions.

In 2021, Fijalkow et al. introduced DeepSynth¹, a general-purpose program synthesizer which constructs programs from input output examples [20]. They make use of a 2-step pipeline in which a neural network learns to predict weights for a context free grammar (CFG) and then a search algorithm seeks programs matching the query. Two main strategies are introduced: enumerative search and sampling.

There is a trade-off between finding many programs rapidly, or fewer programs that are more likely to solve the problem.

In their framework, we are given a list of primitives, the initial DSL, with a number of syntactic constraints

which compile into a context free grammar. they then train a prediction model to predict the weights of the cfg making it a pcfg. lastly they search through the pcfg to find programs meeting the given specification.

In this work, I am implementing a novel program synthesizer, which integrates with the DeepSynth framework.

I am separating the generative model (world model) from the inference machine. And am using a full transformer. Tasks are encoded using the transformer encoder. [explain in detail], which states are encoded using the decoder. A state is represented as a rule of the cfg. at each step of the trajectory, the combined inputs are used for the forward logit network to predict logits over possible actions which are then added to the state.

9.1 Complexity analysis

¹<https://github.com/nathanael-fijalkow/DeepSynth>

References

- [1] Joshua S. Rule, Joshua B. Tenenbaum, and Steven T. Piantadosi. The Child as Hacker. *Trends in Cognitive Sciences*, 24(11):900–915, November 2020.
- [2] Kevin Ellis, Catherine Wong, Maxwell Nye, Mathias Sablé-Meyer, Lucas Morales, Luke Hewitt, Luc Cary, Armando Solar-Lezama, and Joshua B. Tenenbaum. DreamCoder: bootstrapping inductive program synthesis with wake-sleep library learning. In *Proceedings of the 42nd ACM SIGPLAN International Conference on Programming Language Design and Implementation*, pages 835–850, Virtual Canada, June 2021. ACM.
- [3] Stanislas Dehaene, Fosca Al Roumi, Yair Lakretz, Samuel Planton, and Mathias Sablé-Meyer. Symbols and mental programs: a hypothesis about human singularity. *Trends in Cognitive Sciences*, 26(9):751–766, September 2022.
- [4] Volker Peckhaus. Leibniz’s Influence on 19th Century Logic. In Edward N. Zalta, editor, *The Stanford Encyclopedia of Philosophy*. Metaphysics Research Lab, Stanford University, Winter 2018 edition, 2018.
- [5] Michael Rescorla. The Language of Thought Hypothesis. In Edward N. Zalta, editor, *The Stanford Encyclopedia of Philosophy*. Metaphysics Research Lab, Stanford University, Summer 2019 edition, 2019.
- [6] Panu Raatikainen. Gödel’s Incompleteness Theorems. In Edward N. Zalta, editor, *The Stanford Encyclopedia of Philosophy*. Metaphysics Research Lab, Stanford University, Spring 2022 edition, 2022.
- [7] B. J. Copeland. *The Essential Turing*. Oxford University Press UK, 2004.
- [8] Noam Chomsky. On certain formal properties of grammars. *Information and control*, 2(2):137–167, 1959.
- [9] Fosca Al Roumi, Sébastien Marti, Liping Wang, Marie Amalric, and Stanislas Dehaene. Mental compression of spatial sequences in human working memory using numerical and geometrical primitives. *Neuron*, 109(16):2627–2639.e4, August 2021.
- [10] Steven T Piantadosi. The computational origin of representation. *Minds and machines*, 31(1):1–58, 2021.
- [11] Steven T Piantadosi and Felix Hill. Meaning without reference in large language models. *arXiv preprint arXiv:2208.02957*, 2022.
- [12] Quan Do and Michael E Hasselmo. Neural circuits and symbolic processing. *Neurobiology of Learning and Memory*, 186:107552, 2021.
- [13] Adam Santoro, Andrew Lampinen, Kory Mathewson, Timothy Lillicrap, and David Raposo. Symbolic behaviour in artificial intelligence. *arXiv preprint arXiv:2102.03406*, 2021.
- [14] Douglas R. Hofstadter. *Gödel, Escher, Bach: an Eternal Golden Braid*. Basic Books Inc., 1979.
- [15] Artur d’Avila Garcez and Luis C Lamb. Neurosymbolic ai: the 3rd wave. *arXiv preprint arXiv:2012.05876*, 2020.
- [16] Geoffrey E Hinton, Peter Dayan, Brendan J Frey, and Radford M Neal. The “wake-sleep” algorithm for unsupervised neural networks. *Science*, 268(5214):1158–1161, 1995.
- [17] Emmanuel Bengio, Moksh Jain, Maksym Korablyov, Doina Precup, and Yoshua Bengio. Flow Network based Generative Models for Non-Iterative Diverse Candidate Generation. In *Advances in Neural Information Processing Systems*, volume 34, pages 27381–27394. Curran Associates, Inc., 2021.
- [18] Anirudh Goyal and Yoshua Bengio. Inductive biases for deep learning of higher-level cognition. *Proceedings of the Royal Society A: Mathematical, Physical and Engineering Sciences*, 478(2266):20210068, October 2022. Publisher: Royal Society.
- [19] Yoshua Bengio. Scaling in the service of reasoning & model-based ML. <https://yoshuabengio.org/2023/03/21/scaling-in-the-service-of-reasoning-model-based-ml/>, 2022. Accessed 21-08-2023.
- [20] Nathanaël Fijalkow, Guillaume Lagarde, Théo Matricon, Kevin Ellis, Pierre Ohlmann, and Akarsh Potta. Scaling Neural Program Synthesis with Distribution-based Search, October 2021. arXiv:2110.12485 [cs].